

When Does Soccer Need Mixed Strategies?

A pure-first Nash-Q approach: where
a two-player soccer Markov game
does — and does not — need mixing,
and how much LP work that saves.


FOUR QUESTIONS

RQ1 — when does the soccer Markov game admit pure versus mixed stage equilibria?

RQ2 — can pure-saddle detection reduce Nash-Q computation without changing the exact value? **RQ3** — what spatial and transition structures characterize the mixed-equilibrium states? **RQ4** — how sensitive are classification and computation to discounting, tolerance, and reward formulation?

repo · soccer-markov-nash | 315+ tests | game family in docs/design.md | methods
in docs/methods.md | every policy drawn in docs/gallery.html

Two documents, by design: this is the technical notebook — for the plain-language companion, see [The Mixed Game](#) (neither is a stub of the other).



The kickoff. 7×5 grid, 2380 non-terminal states. Player 0 (blue) carries the ball and attacks the right edge; player 1 (green) the left. Start positions are the ones the A10 page assigns this netID. Actions U D L R are chosen simultaneously; reward is +1 / -1 / 0, zero-sum.

POSITION RELATIVE TO LITTMAN (1994)

Littman **already established** that the soccer Markov game can require mixed strategies — his Figure 2 gives a state where any deterministic offensive choice is blocked indefinitely and only randomising between stand and move guarantees an opening. This project does not rediscover that. It asks: **where does mixing appear, why does it appear, and what parameter switches it on?** Littman demonstrated the phenomenon in a *particular* soccer Markov game; we investigate which structural features determine when mixed equilibria arise, and whether their computational cost can be avoided except where mixing is genuinely necessary (docs/littman.md reproduces his Figure 2 and Table 3).

ABSTRACT

- 1 **The soccer game returns pure strategies until the collision rule couples the ball-steal to both players' actions.** Deterministic or coin-flip resolution $\rightarrow$ a pure saddle at every one of the 238 000 (state $\times$ step) stage games (also positionally determined). Littman's random *move order* $\rightarrow$ mixing, but only when it is the majority rule (sharp threshold at $\text{blend} \approx 0.5$ in this interpolation family) and only with a multi-cell goal. That last clause is specific to the move-order rule: action-independent movement noise (slip) forces mixing at any goal width.
- 2 **The pure-first hybrid eliminates LP work on almost every stage game** — 100% of states on the deterministic game, 96.05% on the random one — matching the all-LP value function to $4\text{e-}16$. That is a property of the game; the wall-clock speedup is reported separately.
- 3 **Structural characterization with a per-stage-game certificate:** the mixed states combine stochastic move order, a multi-cell goal mouth, and a defender with insufficient one-step coverage; *every* tested mixed stage game contains a 2×2 matching-pennies submatrix, and the strongly-mixed core is ~ 26 (entropy ≥ 0.9 bits). The single-cell $\rightarrow$ pure / wider $\rightarrow$ mixed switch is machine-certified invariant to the discount, the action set, and the reward objective (make `verify`); the region carries 41% of the equilibrium-path occupancy at 4% of the state space.
- 4 **Action accuracy is not equilibrium accuracy.** A policy network that names the exact-optimal action 99.9% of the time is more exploitable than a mediocre value network. On the harder game where mixed equilibria genuinely matter, it gets sharper still: fitting a Q-net directly to the exact stage matrices is the closest numerical fit and the single most exploitable policy of the four baselines tested — matching a continuous target well and matching a discrete equilibrium well are different objectives. The exact solver stays the ground truth.

READING GUIDE: CORE, SUPPORTING, EXPLORATORY

This report covers more ground than its central claim needs (full accounting in §11, “What weight each experiment carries”). If time is short:

- **Core** — the pure-first hybrid algorithm and its LP-elimination result (Method, RQ2); where and why mixing occurs (RQ3, Mechanism); and **the practical payoff: a minimax policy exploits weak opponents and survives a challenger built specifically to beat it, while every deterministic policy — including a hand-scripted one that beats a random opponent ~76% of the time — does one or the other, never both** (“Why it matters” below). That last result is the direct answer to *why any of this matters*: a predictable pure policy can be learned and exploited, and the 94 no-pure-saddle states are exactly where that predictability would show up.
- **Supporting validation** — numerical robustness across discount and tolerance (RQ4); self-play and occupancy on the equilibrium path (docs/occupancy.md); geometric templates and symmetry (docs/templates.md); which no-pure-saddle states are a uniquely forced mix vs. degenerate (docs/degeneracy.md).
- **Exploratory extensions**, groundwork for future work, not equal-weight contributions — neural function approximation, including on the harder random-move-order game; general-sum solving; reward shaping; policy iteration; the single-cell proof machinery; the phase diagram; the transition-rule variants (tackle, slip, blend). The eventual actual challenge — a continuous-action version of this game — is future work in the same sense; see the roadmap at the end of the Exploratory section.

Method

CONTRIBUTIONS AND HYPOTHESIS

1. Algorithmic: a pure-first hybrid Nash-Q backup — check for a pure saddle ($O(A^2)$), use it, LP only where it fails; exact everywhere; eliminates LP work on 96–100% of stage games. **2. Empirical structural characterization** (over the tested parameter grid, not a theorem) of where the mixed region is and what switches it on.

Hypothesis: mixed equilibria are not distributed arbitrarily — they arise from a specific interaction between goal-mouth geometry and stochastic move-order coupling, and because most states keep pure saddles a pure-first solver gets the exact value while substantially cutting LP usage.

OBJECTIVE DEFINITIONS – RELATED BUT NOT INTERCHANGEABLE

WHERE	OBJECTIVE
A10 Part 1 / exact	undiscounted, finite 100-step horizon (run_finite_horizon, non-stationary v)
stationary Nash-Q	discounted infinite-horizon fixed point, $\gamma < 1$ (run()) – a contraction proxy for the above
scoring="rate"	continuing discounted: a goal resets, value = expected discounted goal difference
Littman replication	Littman's discount / draw interpretation, $\gamma = 0.9 \equiv 0.1$ per-step draw probability


Unless a section says otherwise, results are the stationary discounted fixed point at $\gamma = 0.9$; RQ4 shows the pure/mixed classification is stable across $\gamma \in [0.5, 0.995]$ and matches the exact undiscounted answer.

The game is zero-sum, so at every state the stage game is a payoff matrix $M(s) [a_0, a_1] = E[R + \gamma V(s')]$ and the Shapley (1953) fixed point is the minimax value:

$$\begin{aligned} V(s) &= \text{val}(M(s)) = \max_n \min_{k \in \mathcal{K}} (p^T M) [a_1] \\ &= \min_p \max_{a_0} (Mq) [a_0] \end{aligned}$$

The $Q = R + \beta \text{Nash}(Q')$ update is this operator. The A10 game is *undiscounted* (± 1 at a goal, tie after 100 steps); $\gamma < 1$ is a contraction device for value iteration and RQ4 shows the split holds across γ . Three stage backups: **pure**


(the maximin security value — a fast lower bound, not a Nash value when it is below the minimax), **mixed** (the LP minimax value, via scipy HiGHS), and **hybrid** (the pure saddle value where maximin = minimax, the LP elsewhere — exact everywhere). Support enumeration (support_enum.py) finds *all* equilibria, including of general-sum games.



The pure-first hybrid. On the deterministic game the "yes" branch is taken at every one of the 2380 states, so the LP is never called.

RESULTS

Where mixing comes from, and what the solver does about it



The whole story. Left: the mechanism chain – goal width, then stochastic move order and interception geometry, produce crossing best replies and a matching-pennies stage game where the LP is needed. Right: the three backups, the experimental result, and the occupancy headline.

RQ1 – pure versus mixed, by collision rule


The A10 game is undiscounted, so it is solved exactly by 100-step backward induction (run_finite_horizon, experiments/undiscounted.csv):

EXACT UNDISCOUNTED GAME – 238 000 (STATE × STEP) STAGE GAMES

MOVE ORDER	MIXED STAGE GAMES	V(KICKOFF)	PURE EQUILIBRIUM?
deterministic (exact A10)	0	0.000	yes – pure (non-stationary)
coinflip tie-break	0	0.000	yes
random move order	3 148 (404 states)	+0.459	no

The deterministic equilibrium is **constructive and memoryless**: each player's win attractor (states from which it forces a goal, soccer_nash/attractor.py) equals its $V^* = \pm 1$ set, and the pure profile "attractor move on your winning set, safety move elsewhere" realizes V^* at every state (scripts/positional.py) — the game is positionally determined.

carrier position (player 0)



V^* by player 0 position (blue = player 0 ahead)

defender position (player 1)



V^* by player 1 position (blue = player 0 ahead)

The **value surface** of the random-order game ($\gamma = 0.9$). Left: V^* by carrier position – a smooth gradient toward the attacking goal, saturating to +1 on the two cells that score in one move. Right: V^* by defender position with the carrier held near midfield – nearly flat, the defender's location barely moves the value.

The stationary $\gamma < 1$ solve agrees and is faster — 0 no-pure-saddle stage games on the deterministic game at every γ from 0.5 to 0.995, and 94 / 2380 on the random game at $\gamma = 0.9$ ($V(\text{kickoff}) = +0.164$ at the netID start). The no-pure-saddle count never depends on the kickoff.

IT IS THE MOVE ORDER, NOT THE RANDOMNESS – AND IT HAS TO BE THE MAJORITY


The deterministic and coin-flip value functions are *identical* ($\max |V_{\text{det}} - V_{\text{coin}}| = 0$ at every γ): optimal play never enters a losing contest, so the coin is never flipped. Only Littman's random move order — where a ball-steal depends on who resolves first *and* on both players' targets — creates matching-pennies subgames. Interpolating the two (move_order="blend", resolve by random order with probability blend) shows the onset is **sharp: 0** mixed stage games for $\text{blend} \leq 0.5$, dozens at $\text{blend} = 0.52$, and an *overshoot* to $\sim 1.5\times$ the fully-random count near $\text{blend} \approx 0.7$ before it relaxes (scripts/blend.py, docs/blend.md).

Scope of this claim: in this particular soccer family, this one coin-flip collision variant does not create new mixed stage games, whereas random move ordering does — and only when it is the majority rule in the $P = (1-p)P_{\text{det}} + pP_{\text{random}}$ interpolation. This is *not* a statement about all action-independent stochastic transitions, which can certainly alter equilibrium structure in other games. Whether $p_c \approx 0.5$ is structural is open.




A threshold, not a ramp. No-pure-saddle stage games vs. the fraction of moves resolved by random order (5x4 board). Nothing until the random rule is the majority; then a jump and an overshoot.

random order -> mixed



mixed equilibrium · $V = +0.078$

deterministic -> pure



pure equilibrium · $V = +0.000$

One state, two resolution rules. Arrows are the equilibrium action distribution, sized by probability. Random order (Littman) makes the stage game matching pennies; deterministic carrier-wins collapses it to a pure best reply. Full set: [docs/gallery.html](https://docs.google.com/gallery).

RQ2 – Efficiency

Two numbers, kept apart: the **LP-call rate** is a property of the game and reproduces exactly; **wall clock** is a median of 5 repeats and depends on the machine and LP backend (`scripts/benchmark.py`).

RANDOM MOVE ORDER, $\Gamma = 0.9$ – THE CASE WHERE THE LP IS ACTUALLY NEEDED

SOLVER	SWEEPS	LP CALLS	STATES NEEDING LP	MEDIAN WALL CLOCK
pure (lower bound)	18	0	–	0.3 s
hybrid exact	108	9 672	3.95%	13.2 s
all-LP exact	108	~3.6 M	100%	~8 min

The portable claim: the hybrid eliminates LP work on almost every stage game — 100% of states on the deterministic game, **96.05%** on the random one — and matches the all-LP value to $4e-16$. That is a property of the game, not the machine. The observed wall-clock ratio (~13 s vs ~8 min) is reported separately: it also reflects matrix construction and the LP backend. **Mirror symmetry** (`run_symmetric`, deterministic dynamics only): every state has a distinct board-flip-plus-player-swap mirror, so the 2380 states are 1190 pairs — reconstruct one from the other by $V(\text{mirror}(s)) = -V(s)$ (and the equilibrium *policies* are equivariant too, where unique: `verify_policy_equivariance` returns 0 for the random game's 56 mixed states).

3.95% OF STATES, BUT 41% OF THE EQUILIBRIUM PATH


The LP-call rate counts the whole state space. Under the Nash policy from the kickoff only **456 of 2380** states are ever reached, and the 94 no-pure-saddle states carry **41%** of the discounted occupancy — a 2× over-representation on the path, 10× vs. the raw count (`soccer_nash/occupancy.py`, `docs/occupancy.md`). The carrier's optimal march to goal runs straight through the contested band; four of the eight most-visited states are mixed. The mixed structure is rare to *store* and common to *use*.

Freeze-then-iterate, and its staleness

RANDOM MOVE ORDER, HYBRID SOLVER – SAME FIXED POINT, $|\text{DIFF}| < 3E-9$

	ROUNDS / SWEEPS	MATRIX-GAME SOLVES	WALL CLOCK
value iteration	108	9 672	13.8 s
policy iteration	21	1 879	17.2 s


5× fewer LP solves — but the frozen strategies' distance to the current Nash stays *maximal* (a full pure-strategy flip) for **16 outer rounds**, then collapses to $< 1e-8$. The thrashing eats the saving; on the deterministic game it is strictly worse. Value iteration with the per-sweep Nash cache wins here.



Value iteration converges cleanly. The Bellman residual decays geometrically at rate $\approx \gamma - 108$ sweeps to $1e-9$. Freeze-then-iterate, by contrast, makes *no* progress until the frozen stage saddle flips (next figure) — concurrent stochastic games have no monotone policy improvement (docs/discussion.md §4).

WHICH VALUE QUANTITY TO BACK UP IN THE FROZEN SWEEP

A specific open question: the frozen (p, q) don't match the current Q , so does the evaluation sweep back up $p^T M q$, $\min_k (p M)_k$, or $\max_i (M q)_i$? All three reach the *same* fixed point (zero-sum equilibria are interchangeable), but **$p^T M q$ converges in half the rounds** (21 vs 41–42) — it is the unbiased estimate while the strategies are stale, where the two bounds are systematically pessimistic / optimistic and stretch the thrash phase from 16 rounds to 28 (run_policy_iteration(eval_value=...), docs/numerics.md §3).




Back up $p^T M q$. Staleness of the frozen strategies per outer round. The middle quantity settles by round 16; the pessimistic $\min_k (pM)_k$ and optimistic $\max_i (Mq)_i$ keep flipping until round 28.

RQ3 – what characterizes the mixed states


Across the tested family, a one-cell goal produced no mixed stage games while every tested goal width ≥ 2 produced some. The phase diagram sweeps 127 configurations — boards with $3 \leq \text{width} \leq 11$, $3 \leq \text{height} \leq 9$, $\text{width} \cdot \text{height} \leq 45$, against goal widths 1 to height-2 (scripts/phase_diagram.py --analyze). Precisely: a multi-cell goal is *necessary* for the mixed region in the studied family, and across the tested boards *sufficient* for at least one mixed state — not a proof that every such board must have one, nor that every state is mixed.

	goal-mouth width						
	1	2	3	4	5	6	7
3x3	0						
3x4	0	12					
3x5	0	10	9				
3x6	0	7	7	7			
3x7	0	7	6	6	6		
3x8	0	6	6	5	5	5	
3x9	0	6	5	5	4	4	4
4x3	0						
4x4	0	8					
4x5	0	8	6				
4x6	0	7	6	5			
4x7	0	7	5	5	4		
4x8	0	6	5	5	4	4	
4x9	0	6	5	4	4	4	3
5x3	0						
5x4	0	7					
5x5	0	6	4				
5x6	0	6	4	4			
5x7	0	5	4	4	3		
5x8	0	5	4	3	3	3	
5x9	0	5	4	3	3	3	3
6x3	0						
6x4	0	8					
6x5	0	6	4				
6x6	0	5	4	3			
6x7	0	4	4	4	3		
7x3	0						
7x4	0	7					




The gate (127 configs, $3 \leq W \leq 11$, $3 \leq H \leq 9$, $W \cdot H \leq 45$). Goal width 1 $\rightarrow$ 0 mixed stage games on all 40 tested one-cell configs. Goal width $\geq 2 \rightarrow$ at least one on all 87 tested wider-goal configs. The *fraction* (2.7–12%) is an OLS R^2 0.64 in board area – a dilution effect – goal width adding little with a negative coefficient.

defender at (5, 1)




cells: minimax minus maximin of the carrier's stage game

defender at (3, 2)



cells: minimax minus maximin of the carrier's stage game

defender at (5, 2)



cells: minimax minus maximin of the carrier's stage game

Where the carrier has to guess. Defender pinned at the green disc; each cell shaded by minimax – maximin of the carrier's stage game (0 = pure saddle). Mixing is a local phenomenon – a handful of cells where the defender contests the forward lane.


Geometry predicts the label (scripts/geometry_model.py): a depth-4 decision tree on carrier-frame features separates mixed from the rest with **precision 0.96, recall 0.91**. The dominant feature is defender_can_intercept — the defender within one move of the carrier's forward cell ($P(\text{mixed}) = 0.44$ vs 0.002).

Beyond prediction, the 94 no-pure-saddle states canonicalize under the board mirror to 47 pairs and cluster into **8 geometric templates** (scripts/templates.py, docs/templates.md).

THE MIXED-STATE TEMPLATES – CARRIER-FRAME GEOMETRY, EQUILIBRIUM SUPPORT

STATES	GEOMETRY	SUPPORT	STRUCTURE
24	defender 1 ahead, carrier one row off the goal row	2×2	matching pennies
24	defender 1 ahead, same row, carrier on a goal row	2×2	matching pennies
16	defender 2 ahead, same row, carrier on a goal row	2×2	matching pennies
4	defender diagonally adjacent, carrier off the goal row	2×2	matching pennies
14	defender 1–2 ahead, same row, on a goal row	2×1	near-pure saddle
8	defender 2 ahead, same row, on a goal row	3×2	degenerate (row ties)
4	defender 2 ahead, carrier off the goal row	3×3	wider mix

The four 2×2-support templates are all verified matching pennies — the carrier's best reply flips between the two defender columns and vice versa — covering 68 of 94 states: the carrier choosing between two of {U, D, L, R} against the defender covering one of the same two (exact letters per template in docs/templates.md). All 94 share one value across every equilibrium; 64 have a unique one.



One policy fan per matching-pennies template. Each covers a class of mixed states related by carrier-frame geometry; p^* , q^* are the equilibrium weights on the first support action. The carrier randomizes between two scoring lanes, the defender between covering them.

The mechanism (machine-certified per config; interpretable, not a general theorem): a stage game requires mixing when a **stochastic ($\frac{1}{2}$ - $\frac{1}{2}$) resolution order** meets a **multi-cell goal** and a **defender that can contest the forward cell but not cover both lanes**. Then the carrier's "which lane" reply flips with the defender's cover and vice versa — *crossing best replies*, hence no pure saddle:


geometric condition $\Rightarrow$ crossing best replies $\Rightarrow$ no pure saddle

Drop any one clause — a one-cell goal, deterministic resolution, the coin-flip tie-break — and every tested stage game has a pure saddle. docs/mechanism.md gives the interpretable argument for the single-cell case. What is *not* settled is a board-size-free theorem for arbitrary W , H . *Supporting theory only* (per the meeting feedback, not a deliverable): for a single goal cell the defender's optimal strategy is closed-form and machine-verified, and every single-cell stage game is dominance-solvable, checked for all boards up to 11×5 .

A PER-STAGE-GAME CERTIFICATE

soccer_nash/certificate.py re-derives the pure/mixed label from the stage matrix with $O(A^2)$ arithmetic and no LP; scripts/verify_mechanism.py (make verify) runs it over the reachable state space. Single goal cell $\rightarrow$ a saddle-cell certificate at every state (0 mixed, cross-checked against onecell.certify). Wider goal $\rightarrow$ every mixed stage game gets a certificate an independent checker re-verifies to $\sim 1e-16$: the maximin $<$ minimax gap (the LP-free proof that no pure saddle exists), a closed best-reply cycle, the exact equilibrium ($\epsilon < 1e-15$), and a 2×2 **matching-pennies submatrix** — **present in all 94 of the 7×5 no-pure-saddle states and all 56 of the 5×4 ones**, strictly stronger than the 68-of-94 support count (the core is a property of the matrix; the LP's *reported* equilibrium can still be degenerate on 30 of the 94 without changing whether the core is there — docs/degeneracy.md). The switch is checked invariant across $\gamma \in \{0.5, 0.9, 0.995\}$, 4 vs 5 actions, win vs rate scoring, and exact undiscounted backward induction (16/16), and a $\pm 1e-6$ perturbation flips no genuine classification. docs/result.md is the precise statement and the prior-work positioning.

Mixed stage game: the 2×2 matching-pennies core




The 2×2 matching-pennies core of a mixed stage game. Blue bar = player 0's best row per column; green bar = player 1's best column per row. They never coincide — the best replies cycle, so mixing is forced. Every tested mixed stage game contains one.

The same fact, drawn as a node-and-arrow graph instead of a highlighted grid (one lit-up cell for a pure state, a closed loop of arrows for a mixed one) next to where the players actually are on the board, across twelve representative cases with the full 4×4 matrix printed for each — docs/positions.md / docs/positions.pdf, make positions.


How far the goal-width switch generalizes (scripts/generalize.py, docs/generalize.md). It survives board scale (checked to 11×5 and 9×7), aspect ratio, and goal placement: non-contiguous goal cells like (1, 3) or (0, 2, 4) still produce mixed stage games, so adjacency does not matter, only that the carrier has two cells the defender cannot cover at once. It does *not* survive a change of transition family. Adding action-independent movement noise (SoccerGame(slip=p): each player takes a random move with probability p) produces mixed stage games even under deterministic resolution and even with a single goal cell — a 5×5 one-cell board goes from 0 mixed to 62 at slip = 0.1. The same is true of this project's own tackle rule — the defender commits to a challenge that wins the ball with probability tackle_prob or bounces off (docs/tackle.md). So the switch is a property of Littman's move-order rule, where the unresolved coin only bites where the carrier has two lanes; slip and tackle supply that coin everywhere, a fair coin flip on contested squares supplies it nowhere on the

equilibrium path. Reading the 2×2 cores of the 7×5 mixed states, the carrier's crossing pair is a vertical move in 90 of 94 — it is choosing which goal row to head for, and the defender is guessing it.


deterministic




coinflip




random




blend 0.7



slip 0.2



tackle 0.5



A fingerprint per collision rule. Same board, same pinned defender; shading is $\text{minimax} - \text{maximin}$ of the carrier's stage game. Each rule for resolving a contested cell leaves its own mark: the deterministic and coin-flip rules stay pure everywhere, Littman's move order and its blend light a thin band at the goal mouth, and the two rules that put an unavoidable coin on every square — slip and tackle — light a broad region around the defender.

EXISTENCE IS NOT THE SAME AS AMOUNT

Every one of the 94 no-pure-saddle stage games sits within 0.07 of a pure saddle ($\text{minimax} - \text{maximin}$, median 0.029), and the carrier's **mixing entropy** has median 0.55 bits — most are near-pure hedges ($\approx 87/13$), not real randomisation. Only 26 of 94 reach ≥ 0.9 bits (a near-even split); those are the states with genuine indifference. So the mixed region is 94 states, but the *strongly-mixed* core is ~ 26 . Robust structure: 68 have a 2×2 matching-pennies support; the *value of mixing* ($V(\text{hybrid}) - V(\text{pure maximin})$) averages +0.13 there and compounds — at the centred kickoff a pure player secures only a draw where mixing is worth +0.15 (scripts/mixing.py).


The reward objective does not move the mixed region

The default game rewards a *win* — first goal ends it, value $P(\text{win}) - P(\text{loss})$. A second objective, `scoring="rate"` (`scripts/reward.py`), keeps playing: a goal scores ± 1 and restarts with the ball to the conceding team, so the value is the expected discounted *goal difference* and $\gamma < 1$ is load-bearing. On the 7×5 random-order board:

SAME DYNAMICS, TWO REWARD OBJECTIVES – RANDOM MOVE ORDER, DISCOUNT 0.9

OBJECTIVE	NO-SADDLE STATES	VALUE RANGE	V(KICKOFF)
win – $P(\text{win}) - P(\text{loss})$	94	$[-1.00, +1.00]$	+0.150
rate – expected goal difference	94 – the same 94	$[-0.88, +0.88]$	+0.132

The no-pure-saddle set is *identical* — 0 states enter or leave. A mixed Nash equilibrium is an equilibrium of a single stage game $M(s)$; changing how the horizon is scored rescales $M(s)$ through $V(s')$ but does not cross the `maximin = minimax` boundary. The value *range* compresses ($\pm 1 \rightarrow \pm 0.88$): under `rate` a lost position is not a cliff — conceding restarts play with possession, a floor under every state. Details in `docs/reward.md`.




Same shape, compressed range. `V*` by carrier position under the two objectives. `rate` pulls the goal-adjacent `+1` cells down to `+0.88` — the concede-and-restart floor — without changing where the value gradient runs.

Littman's fifth action

Littman's 1994 soccer game has a fifth action, `stand`, and his Figure 2 — a carrier pinned against its own goal by an adjacent defender — is the textbook "soccer needs mixing" example: the carrier must randomize between standing and moving. Adding `stand` to the random-order game (`n_actions=5`, `scripts/littman.py`) on Littman's 5×4 board:

ACTION SET	NO-SADDLE STATES	STAND IN THE MIXED SUPPORT	V(KICKOFF)
N, S, E, W	56	0	+0.147
N, S, E, W, stand	56	48	+0.172

The fifth action does *not* move where mixing happens — 48 of the 56 no-pure-saddle states are shared, 8 swap in and 8 out at the margin — but it changes *what* the mix is: the carrier now hedges with STAND in 48 of 56 states rather than with a move action, and every player is slightly better for the option ($V(\text{kickoff}) +0.147 \rightarrow +0.172$). The Figure 2 state (0,1,1,1,1) is a clean matching-pennies matrix with equilibrium ($\frac{1}{2}$ U, $\frac{1}{2}$ STAND) for both players. The deterministic game keeps a pure saddle at every state with five actions too.



Littman's Figure 2. Same state, four moves (left) vs. five (right). The dashed ring is STAND. Adding it turns the carrier's U/L hedge into a U/STAND hedge — Littman's mixed policy.


SO WHAT?

Knowing *where* the 94 no-pure-saddle states are and *why* they arise (above) does not by itself say why a reader should care. The "Why it matters" section below answers that directly: a policy that plays a fixed, predictable action at those states can be learned and exploited by an opponent built specifically to punish it — greedy beats a random opponent by more than minimax does, then loses to its own challenger; minimax beats the same challenger every time. That result is moved up into the reading guide near the top of this report for exactly this reason; RQ4 below is the numerical due-diligence behind the *where*, not the *why it matters*.

RQ4 – Numerical robustness

The soccer game contains rock-paper-scissors

The pure/mixed check is an $O(A^2)$ best-reply method: mark player 0's best row in each column and player 1's best column in each row; a cell with both is a pure saddle, and if none has both the best replies *cycle* and the game needs mixing — no LP required.



Same structure. Left: rock-paper-scissors. Right: the stage game at $(0, 1, 1, 1, 1)$ — carrier pinned against its own goal, defender adjacent (Littman's Figure 2 geometry). Blue bar = player 0's best row per column; green bar = player 1's best column per row. They cross both ways — matching pennies, no pure equilibrium.

The value is three numbers

`value_bracket` returns what the row player guarantees ($\min_k (pM)_k$), gets ($p^T Mq$), and can reach ($\max_i (Mq)_i$) — three quantities that are not guaranteed to be numerically identical. The true value is bracketed, so the width certifies the LP error.


FINDING

With scipy's HiGHS the three agree to $1.1e-16$ per stage game, $8.7e-10$ accumulated — the definitional ambiguity is not a practical problem. If it were (an older vertex solver, or accumulated drift), the row player should take the *guaranteed* value $\min_k (pM)_k$: it is always a safe lower bound on the true minimax.

Discounting and tolerance

	DETERMINISTIC	RANDOM (7×5)
mixed states, $\gamma = 0.5 \rightarrow 0.9 \rightarrow 0.995$	0 → 0 → 0	122 → 94 → 102
classification split, rel_tol 1e-12 → 1e-3	stable	604 / 1682 / 94 (stable)
classification, rel_tol 1e-2 / 1e-1	—	80 / 0 mixed
fixed 0.1-rounding flips a saddle	0	62

Discounting mildly shifts which states mix; the classification is otherwise robust across nine orders of magnitude of tolerance, breaking only as the tolerance approaches the real minimax – maximin gaps. Fixed 0.1-rounding — proposed to force indifference — misclassifies 62 states, exactly the small-entry mixed region.



There is no safe discount. The genuine minimax – maximin gaps (green) never reach the 0.1 rounding grid at any γ : heavy discounting shrinks the γ^* value differences, light discounting leaves entries barely apart. Fixed 0.1-rounding wrongly collapses 60–100% of the mixed games (blue). The scale-aware rel_tol classifier is stable across the whole range.

Why it matters: minimax survives what every deterministic policy cannot


Littman's Table 3: minimax exploits *and* survives

Littman's central experiment trains four policies and scores each against a random opponent, a hand-built one ("simple rules for scoring and blocking", §6.2), and a **challenger** trained to beat it — his minimax policies held ~35–40% of games against their challenger, his Q-learning (deterministic) policies won **0%**. Reproduced here with the *exact* solver on his 5×4 / 5-action / random-order / goal-reset board (scripts/tournament.py, score = player-0 expected discounted goal difference; the hand-built opponent beats a random one ~76% of games):

LITTMAN'S BOARD, $\Gamma = 0.9$ – PLAYER-0 GOAL DIFFERENCE FROM THE KICKOFF

PLAYER-0 POLICY	VS. RANDOM	VS. HAND-BUILT	VS. ITS CHALLENGER	ROBUSTNESS GAP
minimax (Nash)	+0.67	+0.28	+0.15	+0.52
greedy vs. random	+0.92	+0.68	-0.59	+1.51
greedy self-play	+0.05	0.00	-0.00	+0.05
hand-built	+0.72	0.00	-0.68	+1.39

Minimax holds up; greedy collapses to its challenger



Only minimax is in the top-left. greedy/rand is the strongest policy against weak opponents and among the weakest against a challenger – the rock-paper-scissors trap. minimax presses its advantage (+0.67) and still wins (+0.15) against the worst case, because the Nash policy has no exploitable row.

The mechanism is exactly Littman's "every deterministic offense has a perfect defense": the mixed states are matching-pennies stage games; the greedy *and* the hand-built policy each commit to one row of each, and the challenger plays the column that beats it. The minimax policy randomizes with the weights that make the challenger indifferent. Full table (including the $P(\text{win}) - P(\text{loss})$ view) in docs/tournament.md.

docs/tournament_deepdive.md turns that mechanism into a measured, causal claim instead of an assertion: patching greedy's policy with the exact Nash mix at *only* the 7.4% of states that are genuinely mixed — leaving every other state as greedy already plays it — recovers 43% of the gap to minimax against a freshly built challenger, with one concrete exploited state shown matrix and all, and a gamma sweep showing that recovered share shrinks (and briefly reverses) as the horizon lengthens.

Other checks

- Duality gap $V0_{br}(s_0) + V1_{br}(s_0) < 1e-9$ for every move order — no opponent beats the game value.
- Nash vs. Nash reproduces the value: forced draws in the deterministic and coin-flip games; $+0.162 \pm 0.002$ empirical discounted return over 5 seeds of 3000 games (vs. $+0.164$ computed) in the random game.
- A best response to one assumed opponent is fragile: the Part 2 policy has exploitability 0.43 — worse than random — which is the A10 competition's warning about non-equilibrium submissions.
- A10 networks over 5 seeds (experiments/a10_competition_seeds.csv): the Part 2 imitation net plays optimally at 0.990 ± 0.000 of states and wins Q8 every time; competition Network Second is exact, but Network First sits at a 0.19 ± 0.02 exploitability from this kickoff that does not wash out with re-seeding.

Exploratory — function approximation and general-sum

These are groundwork for the eventual continuous-action work, not equal-weight contributions. The exact solver is the point of both: it is the ground truth they are measured against.

Three starting points for the Q net, and where the approximation breaks

On the deterministic game (5 seeds, experiments/nash_dqn_seeds.csv): a **Q net** regresses toward the stage matrices $Q(s, a_0, a_1)$ and extracts a policy by minimax; a **policy net** regresses *directly* onto the exact equilibrium strategies. The Q net is tried from three starting points, same architecture, same scoring: **from zero** (random init, pure TD bootstrap, 600 epochs), **fit to exact** (supervised regression straight onto Q_{exact} , no bootstrap at all, 600 epochs), and **warm start** (the fit-to-exact weights, then the same 600-epoch TD bootstrap as "from zero").

5 SEEDS	EXACT	FROM ZERO	FIT TO EXACT	WARM START	POLICY NET
max $ V - V_{\text{exact}} $	0	0.55 ± 0.02	0.38 ± 0.05	0.49 ± 0.06	—
action agreement	100%	43% ± 2%	45% ± 2%	45% ± 2%	99.9% ± 0.0%
pure/mixed classification	100%	67% ± 2%	47% ± 2%	65% ± 2%	—
max equilibrium regret ϵ_{NE}	0	—	—	—	0.46 ± 0.02
exploitability (duality gap)	<1e-9	0.44 ± 0.05	0.47 ± 0.04	0.38 ± 0.06	0.84 ± 0.00

Fitting the Q net directly to Q_{exact} — no bootstrap, no target network, nothing but supervised regression against the ground truth — still only reaches 45% action agreement, barely above the 43% a random init gets from 600 epochs of TD bootstrap. That is the headline result of the warm-start ablation: the bottleneck is not bootstrap noise. A network of this size, extracting a policy via minimax of its own predicted matrix, cannot represent Q_{exact} closely enough for the *argmax* to survive, even when it is simply told the answer. Continuing training from that fit-to-exact starting point (warm start) leaves action agreement essentially unchanged (45% either way) but does buy back some of what pure supervised fitting lost: pure/mixed classification recovers from 47% back to 65% (against from-zero's 67%), and exploitability drops to 0.38 — the best of the three, beating both from-zero (0.44) and fit-to-exact alone (0.47), in 4 of 5 seeds. So the warm start does not fix the underlying representational problem, but it is not wasted either: the extra bootstrapped training pulls the network back toward respecting the game's saddle structure without giving up the head start on value error.

The policy net names the exact-optimal action at 99.9% of states and is still *more* exploitable than any of the three Q nets (duality gap 0.84 vs 0.38–0.47). Two reasons: the ~0.1% of states where it is wrong are exactly the ones a best-responder attacks (the rock-paper-scissors trap again), and a softmax head necessarily smears what

should be pure strategies, manufacturing exploitable near-indifference. So the failure is not primarily value approximation — it is that "name the right action" is not the same as "play a Nash", and neither is "be told Q_{exact} outright". Even on a small discrete game where the exact solution is a 0.3 s computation, a straightforward neural approximation does not preserve game-theoretic robustness. Keep the exact solver as ground truth.

ACTION ACCURACY IS NOT EQUILIBRIUM ACCURACY

A network that names the textbook-optimal move nearly every time can still be trivially exploited, because the states where it is wrong are not random — they are exactly the ones a best-responder is looking for. Simple single-agent RL that optimizes for "predict the right action" produces a policy that looks excellent on that metric and is not robust in the game-theoretic sense at all; this is the discrete, measured version of that concern.

The harder game: random move order, genuinely mixed equilibria

Everything above runs on `A10SoccerGame` — the **deterministic** game, where every stage game already has a pure saddle. That made the from-zero/fit-to-exact/warm-start comparison a genuine function-approximation study, but not yet an answer to the harder question the review of this section raised directly: *can neural Nash-Q reproduce the exact solver on the game where mixed equilibria actually matter?*

Two real fixes were needed in `soccer_nash/nash_dqn.py`, not just a different game object, since the deterministic-only version was silently wrong on this one: `train_nash_dqn`'s bootstrap target used to be the vectorised *maximin* of the target network's predicted matrix — exact only when every stage game has a pure saddle — and the transition precompute used `game.step`, which *samples* one outcome, exact only when a joint action has a single outcome. Both are now general: `_minimax_batch` is the same pure-fast-path check, vectorised, falling back to the LP only for the predicted-mixed matrices; `game.transitions` supplies the full outcome distribution, averaged exactly. `scripts/nash_dqn_random.py` reruns all four baselines on the same plain 4x4 board as `tournament4.py/tournament_deepdive.py` (5x4, two-cell goals, **random move order**), where 56 of 760 states (7.4%) have no pure saddle — a game the deterministic-game experiment never actually exercised (5 seeds, 600 epochs; `experiments/nash_dqn_random_seeds.csv`):

5 SEEDS	EXACT	FROM ZERO	FIT TO EXACT	WARM START	POLICY NET
action agreement	100%	57% $\pm$ 3%	60% $\pm$ 2%	59% $\pm$ 2%	95.8% $\pm$ 0.4%
exploitability	<1e-9	0.36 $\pm$ 0.04	1.01 $\pm$ 0.27	0.36 $\pm$ 0.08	1.20 $\pm$ 0.01
mean Frobenius error	0	0.42 $\pm$ 0.02	0.32 $\pm$ 0.01	0.39 $\pm$ 0.02	—
mean equilibrium regret	0	0.014 $\pm$ 0.001	0.015 $\pm$ 0.002	0.011 $\pm$ 0.001	—

On this board, fitting directly to Q_{exact} is the best *fit* and the worst *policy*. Fit-to-exact has the lowest matrix error (0.32) and the highest action agreement of the three Q nets (60%) — it is, numerically, the closest to the exact answer — and yet its exploitability (1.01) is roughly 3x *worse* than from-zero's (0.36). Minimizing per-state MSE with no regard for the game's structure can produce a matrix that is uniformly close to correct and still names an exploitable policy, because the states that matter for exploitability are not weighted any differently than the states that don't. TD bootstrapping, even though it fits the matrix values less precisely, produces a policy that is harder to punish — the bootstrap target itself depends on downstream states' *minimax* values, which couples consistency across states in a way independent per-state regression cannot. **Warm-starting recovers almost all of it:** continuing training from the fit-to-exact weights lands at from-zero's exploitability (0.36) while keeping better action agreement (59% vs 57%) and the best regret of the three (0.011) — on this harder game, unlike the deterministic one, warm start is the best of the three Q-net variants outright, not just a partial recovery.

The policy net's own failure is starker here than on the deterministic game. 95.8% action agreement is far better than any Q net's, and its exploitability (1.20) is still the worst of the four — the same "action accuracy is not equilibrium accuracy" lesson, now demonstrated on a game that genuinely requires randomization rather than one where the right answer always happens to be a single action.

Width/depth ablation: is 45% a capacity limit?

Back on the deterministic game (A10SoccerGame, where the warm-start result earlier in this section was also measured): more *training* doesn't move action agreement there — fitting straight to Q_{exact} with zero bootstrap noise lands at the same ~45% as 600 epochs of TD bootstrap from scratch. The remaining question (item 15 of the DQN checklist, "use a sufficiently expressive network") is whether more *capacity* does.

scripts/nash_dqn_ablation.py sweeps width (2 hidden layers, 32–512 wide — extended to 512 specifically because the review this answers said not to limit the network size) and depth (1–4 hidden layers, width 96 fixed) using the fit-to-exact method above (2 seeds/config, 600 epochs; experiments/nash_dqn_ablation.csv):

WIDTH SWEEP (DEPTH 2)	32	64	96	128	192	256	512
params	1.8k	5.6k	11.4k	19.3k	41.3k	71.4k	273.9k
action agreement	43%	47%	45%	45%	46%	45%	47%
exploitability	0.79	0.56	0.47	0.41	0.38	0.25	0.18

DEPTH SWEEP (WIDTH 96)	1	2	3	4	3, WIDTH 256	3, WIDTH 512
params	2.1k	11.4k	20.8k	30.1k	137.2k	536.6k
action agreement	44%	45%	46%	48%	48.1%	47.5%
exploitability	0.82	0.47	0.37	0.25	0.08	0.05

Action agreement plateaus — and stays plateaued all the way to 115x the original parameter count. The best action agreement in the whole grid is depth-3/width-256 (137k params, 48.1%); depth-3/width-512 (537k params, 3.9x more parameters than the best config, 115x the original 5->96->96->16 architecture) does not beat it — it is slightly *lower*, 47.5%. Nowhere in this grid, including its largest member, does capacity buy a path toward the policy net's 99.9%. Item 15 is answered as firmly as this experiment can answer it: this is not an expressiveness problem in the sense of "the network can't fit the numbers, give it more room" — more room stops mattering for this metric long before the grid runs out. **Exploitability and value error, on the other hand, never stop improving with capacity** — exploitability falls monotonically from 0.79 (1.8k params) to 0.05 (537k params, a 15x reduction) with no sign of flattening out even at the largest size tested, and max value error falls from 0.53 to 0.03 over the same range. A bigger net genuinely keeps getting *closer* to Q_exact in the numbers that matter for robustness; it just stops getting closer in a way the discrete argmax reflects, well before it stops getting closer in every other sense. That is the same lesson as the policy net's own failure, from the other direction: matching a continuous target well (small MSE) and matching a discrete decision well (correct argmax) are different objectives, and this architecture is much better at the former than the latter, regardless of size. **Depth beats width at matched budget:** depth-3/width-96 (20.8k params) beats width-128/depth-2 (19.3k params, the nearest matched budget) on every metric — action agreement, exploitability, and value error alike — and depth-4/width-96 (30.1k params) beats width-256/depth-2 (71.4k params, 2.4x the parameters) outright on all three. The single best configuration in the grid for action agreement — depth 3, width 256, 137k params — is still 52 points short of the policy net; the single best for exploitability and value error — depth 3, width 512, 537k params — does not even improve on the 137k-param configuration's action agreement while using 3.9x the parameters to chase the other two metrics further.

Roadmap: this is not the eventual challenge

Everything in this section is discrete: a 4x4 (or 5x5) action grid, a network that outputs a matrix over it. The actual eventual challenge is a **continuous-action** version of this game — a continuous dog-and-sheep / continuous soccer pursuit problem — and the point of that future work is *not* to discretize a continuous action space back down to a grid and reuse everything above unchanged. The intended path:

exact discrete game → neural reproduction of it (this section) → continuous-action stage-game solver

The discrete work above is what a continuous solver will be checked against close to the boundary (fine grids should approach the continuous answer), not a component it reuses directly. This repository deliberately does not contain a half-finished continuous solver: that is future work, kept out of this codebase until it is ready to be done properly rather than bolted onto the discrete `_QNet` architecture above.

General-sum — a sanity check, not a second paper

`markov_game.py` extends pure-first beyond zero-sum: enumerate *all* stage equilibria (`support_enum.py`), select one by "largest sum of values". On Battle of the Sexes this correctly avoids the mixed equilibrium whose value is below either pure one; the soccer game is zero-sum, so it does not touch the main result. Left as future work.

Limitations

- The A10 page redacts the ID-specific start position and goal rows; this repo uses the standard Littman geometry (configurable). Across the phase-diagram sweep every one-cell-goal board has 0 mixed states and every wider-goal board has some; the exact 3.95% is the 7×5 / 3-cell-goal case.
- Several collision sub-cases (carrier vs. stationary opponent, bump = steal, swap possession) are *interpreted* from the two stated A10 rules, not quoted. If course staff intended otherwise, Claim A must be re-measured. See `docs/assumptions.md`.
- **Claim A/B** (a pure equilibrium of the deterministic game) is established constructively — the explicit pure memoryless attractor/safety profile. **Claim C** (the theoretical game necessarily has one, over all rules and settings) is *not*, for the general goal mouth; for the single goal cell it is partly proved (`docs/proof.md`).
- `random` and `coinflip` are different game definitions, isolating the collision rule and the tie-break respectively — not the A10 game.
- The goal-width result is **empirical over the tested grid**: a multi-cell goal is necessary for the mixed region in the studied family and sufficient for at least one mixed state on the tested boards — not a theorem.

What weight each experiment carries

Core. Pure vs. mixed by collision rule (RQ1); the pure-first hybrid, exact, LP eliminated on 96–100% (RQ2); the phase diagram — goal width, not board size (RQ3); geometry → templates → 2×2 matching pennies (RQ3); numerical robustness — the three value numbers, the rounding trap, the classifier (RQ4); Littman Table 3 reproduced and its causal follow-up — minimax exploits *and* survives a tailored challenger, greedy does one or the other, patching greedy with the exact mix at only the no-pure-saddle states recovers 43% of the gap (docs/tournament_deepdive.md) — the practical *why it matters*, not just a validation check.

Strong supporting evidence. Occupancy (4% of states, 41% of the path); the reward-objective comparison (region invariant, value range not); symmetry (value *and* policy equivariance); the blend interpolation; which no-pure-saddle states are a uniquely forced mix vs. degenerate (docs/degeneracy.md).

Future / exploratory. The neural baselines, deterministic and random-move-order games alike; the general-sum extension; freeze-then-iterate and the eval_value choice; the single-cell dominance / attractor machinery.

Related work

- **A. Markov games and minimax-Q.** Littman (1994) — the two-player zero-sum Markov game, minimax by LP, soccer *because* its optimal policy can be probabilistic. The baseline this project builds on.
- **B. Concurrent / reachability games.** de Alfaro–Henzinger–Kupferman on positional determinacy — the frame for the attractor result and the single-cell argument.
- **C. Stochastic games and stationary equilibria.** Shapley (1953); Filar & Vrieze; the orderfield-property literature — the frame for "one target cell vs. many".
- **D. Strategy improvement.** Hoffman–Karp, Condon for turn-based games; the *absence* of a monotone guarantee for concurrent games explains the freeze-then-iterate thrash.
- **E. Reward shaping.** Ng–Harada–Russell potential-based shaping — the frame for the PBRS-leaves-the-equilibrium-exact result.

Prior work: mixed strategies can be necessary in soccer (A); pure-first enumeration before LP (the meeting). *This project adds:* an empirical map of where the mixed region is, what switches it on, how much of the equilibrium path it covers, and a hybrid solver that pays LP cost only there. docs/result.md §6 positions the goal-width switch against Littman (1994), the ordered-field property, concurrent reachability games, and pursuit–evasion in full.

soccer_nash — game.py · a10.py · matrix_games.py · support_enum.py · nash_q.py ·
markov_game.py · numerics.py · dominance.py · attractor.py · geometry.py · symmetry.py ·
tree.py · templates.py · onecell.py · certificate.py · reachability.py · occupancy.py ·
evaluate.py · best_response.py · exploit.py · mlp.py · nash_dqn.py · opponents.py ·
shaping.py · simulate.py · experiment.py · render.py · viz.py

scripts – experiments.py · analyze.py · numerics.py · equilibria.py · geometry_model.py · templates.py · phase_diagram.py · onecell_proof.py · verify_mechanism.py · generalize.py · tackle.py · positional.py · positions.py · benchmark.py · policy_iteration.py · selfplay.py · nash_dqn.py · nash_dqn_ablation.py · nash_dqn_random.py · degeneracy.py · distance_table.py · pure_vs_mixed_exploit.py · render.py · figures_png.py · gallery.py · story.py · showcase.py · littman.py · mixing.py · reward.py · blend.py · occupancy.py · tournament.py · tournament4.py · tournament_deepdive.py · a10_part1.py · a10_part2.py · a10_competition.py

docs – design.md · methods.md · discussion.md · numerics.md · assumptions.md · templates.md · proof.md · geometry.md · mechanism.md · result.md · generalize.md · tackle.md · littman.md · tournament.md · tournament_deepdive.md · positions.md · showcase.md · degeneracy.md · mixing.md · reward.md · blend.md · occupancy.md · findings.md · shaping.md · selfplay.md · a10_part1.md · a10_part2.md · a10_competition.md · review-checklist.md · review-18.md · review-15.md · gallery.html

experiments – baseline.csv · undiscounted.csv · gamma_sweep.csv · tolerance_sweep.csv · board_sweep.csv · goal_mouth_sweep.csv · one_cell_goal.csv · phase_diagram.csv · generalize.csv · blend.csv · tackle.csv · reward.csv · mixing.csv · occupancy.csv · tournament.csv · tournament4.csv · tournament_deepdive.csv · degeneracy.csv · distance_table.csv · pure_vs_mixed_exploit.csv · nash_dqn_seeds.csv · nash_dqn_random_seeds.csv · nash_dqn_ablation.csv · a10_competition_seeds.csv

env: 7×5 grid, 2380 states, zero-sum ± 1 , $\gamma = 0.9$